

Class 7: Machine Learning 1

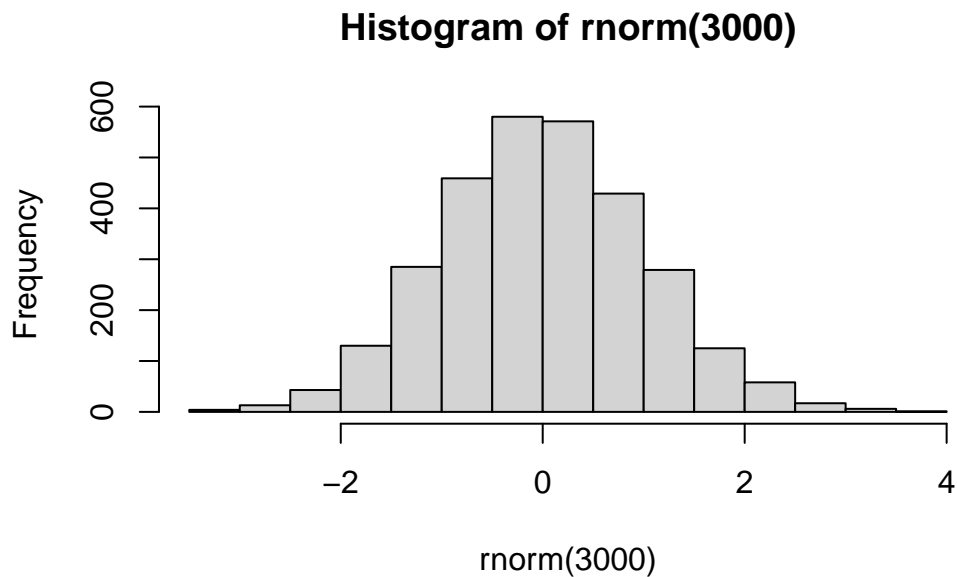
Snehita Vallumchetla (PID A16853399)

Today we will explore unsupervised machine learning methods including clustering and dimensionality reduction methods.

Let's start by making up some data (where we know there are clear groups) that we can use to test out different clustering methods.

We can use the `rnorm()` function to help us here:

```
hist(rnorm(3000))
```



We make data with two “cluster”

```
x <- c(rnorm(30, mean = -3), rnorm(30, mean = +3))

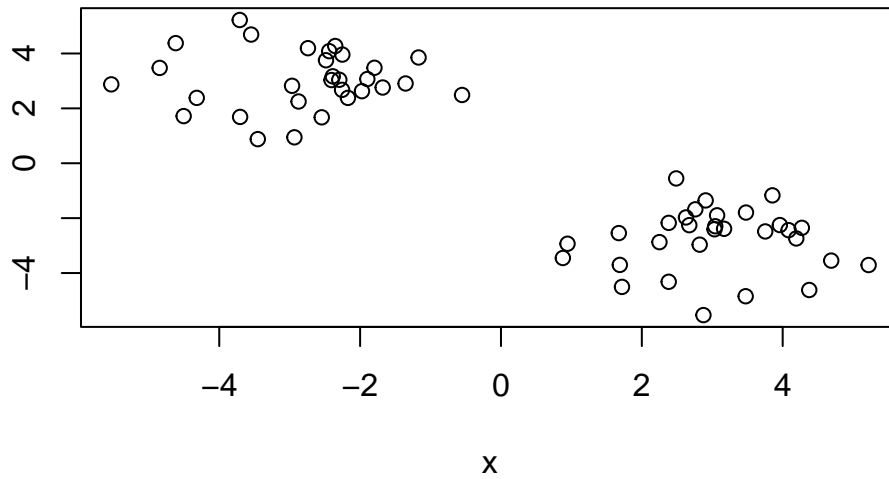
#cbind, causes columns to "bind together" there is also row bind for the binding of rows

z <- cbind(x = x, rev(x))

head(z)
```

```
      x
[1,] -2.173028 2.381752
[2,] -2.873963 2.250486
[3,] -2.966235 2.821831
[4,] -3.709090 5.220407
[5,] -2.545036 1.673360
[6,] -2.740639 4.193083
```

```
plot(z)
```



How big is z

```
nrow(z)
```

[1] 60

```
ncol(z)
```

[1] 2

K-means clustering

The main function in “base” R for K-means clustering is called `kmeans()`

```
k <- kmeans(z, centers = 2)
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

      x
1 -2.792214  3.024492
2  3.024492 -2.792214

```

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 70.8648 70.8648
(between_SS / total_SS = 87.7 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
# shows that each cluster has a size of 30
# clustering vector: it is the membership vector, which cluster do the points belong to
# available components: all of the stuff that is available
```

```
attributes(k)
```

```
$names
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"

$class
[1] "kmeans"
```

Q. How many points lie in each cluster?

```
# size shows how many points are in each vector
k$size
```

[1] 30 30

Q. What component of our results tell us about the cluster membership (i.e. which points lie in which cluster)

```
# shows the cluster that each point is allocated to
k$cluster
```

[illegible]

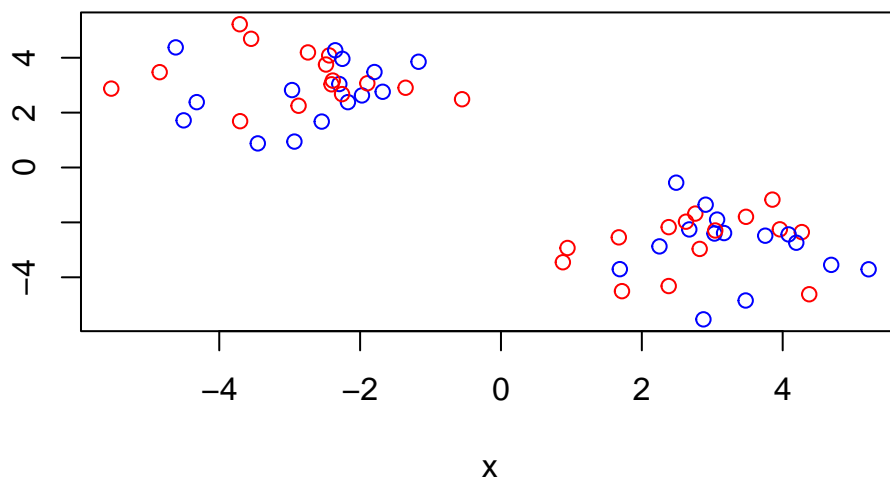
Q. Center of each cluster

```
k$centers
```

	x	
1	-2.792214	3.024492
2	3.024492	-2.792214

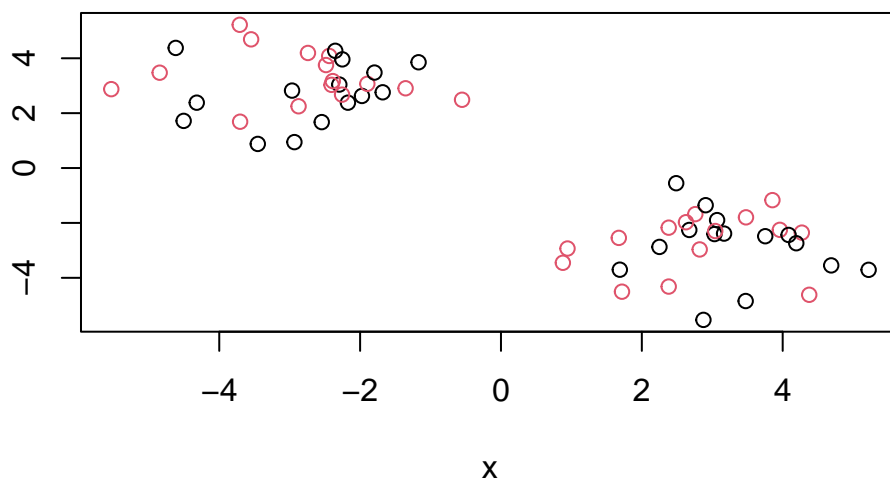
Q. Put this result information together and make a little “base R” plot of our clustering results. Also add the cluster center points to the plot.

```
plot(z, col = c('blue', 'red'))
```

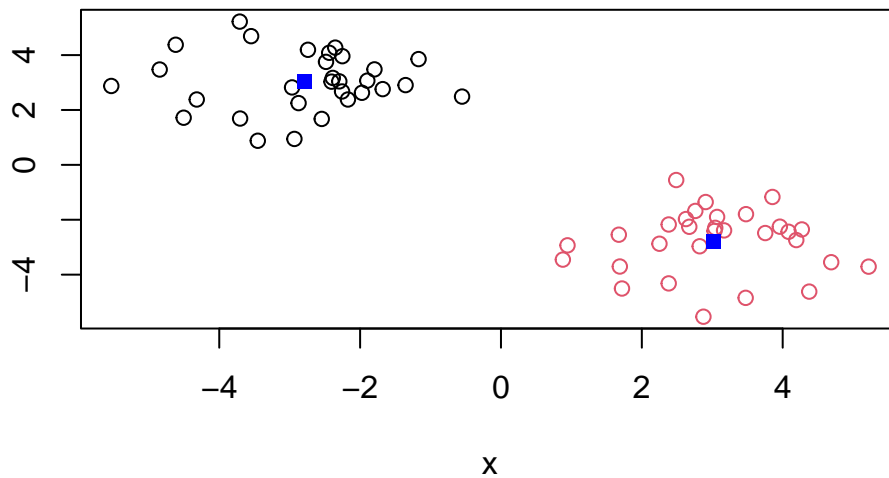


You are able to color plots by the number

```
plot(z, col = c(1, 2))
```



```
# points are now colored by cluster membership
plot(z, col = k$cluster)
points(k$centers, col = 'blue', pch = 15)
```

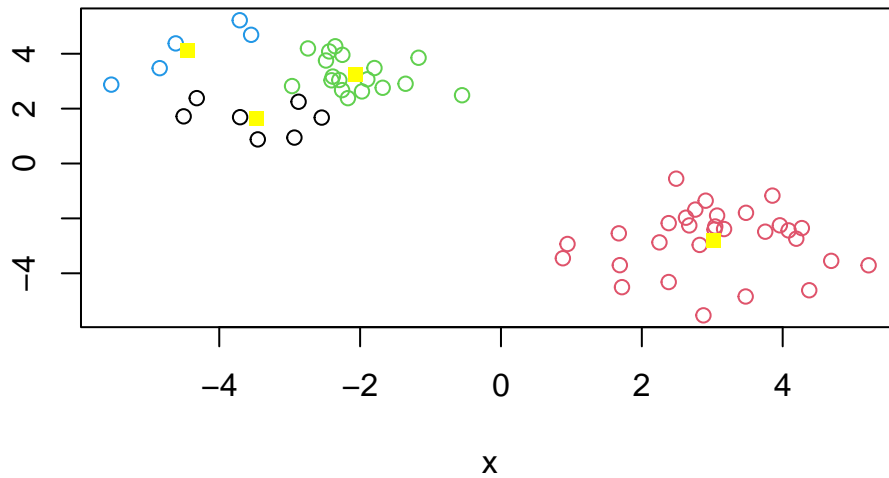


Q. Run kmeans on our input `z` and define 4 clusters making the same result visualization plot as above (plot of `z` colored by cluster membership)

```
k2 = kmeans(z, centers = 4)

plot(z, col = k2$cluster)

points(k2$center, col = 'yellow', pch = 15)
```



Hierarchical Clustering:

The main function in base R for this is called `hclust()` it will take as input a distance matrix (key point is that you can't just give your raw data as input - you have to first calculate a distance matrix from your data).

```
d <- dist(z)

hc <- hclust(d)

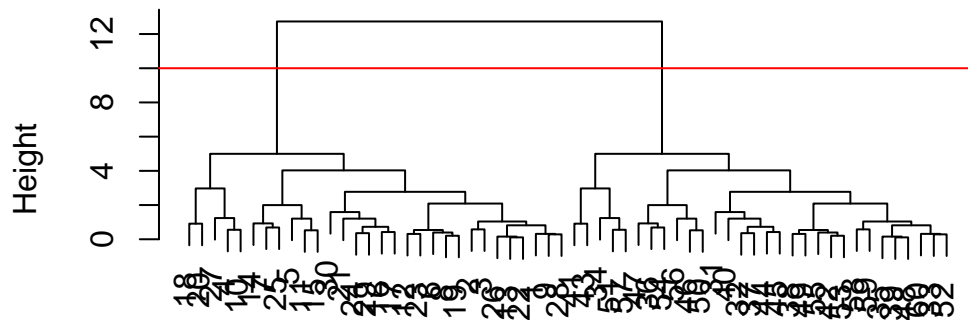
hc
```

```
Call:
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

```
plot(hc)
abline(h = 10, col = 'red')
```

Cluster Dendrogram

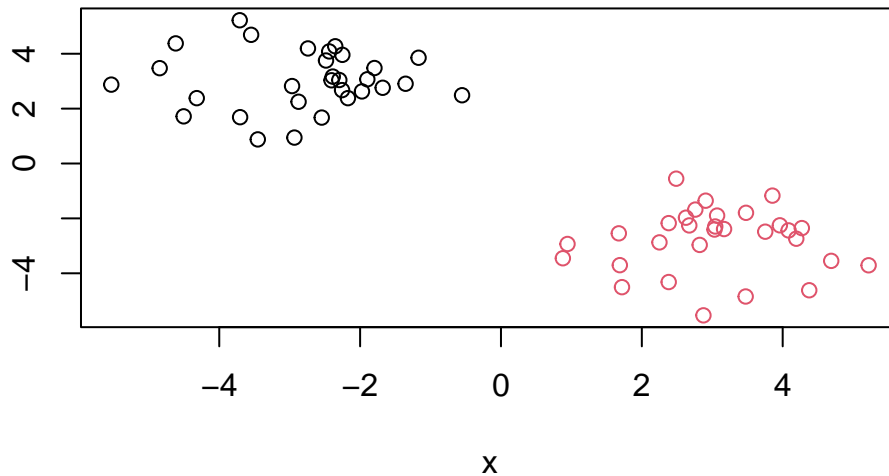


d
hclust (*, "complete")

Once I inspect the “tree” I can “cut” the tree to yield my groupings or clusters. The function to do this is called `cutree()`

```
grps <- cutree(hc, h = 10)
```

```
plot(z, col = grps)
```

Dimensionality Rediction and PCA:

PCA is used to reduce the dimensions in data allowing for something that is able to be visualized. Plot data on a new axis which makes them better for viewing relationships. First principle component describes the most that it can, and the second describes anything else. Filters data, and maintains the main essence of the data.

Let's examine some silly 17-dimensional data detailing food consumption in the UK. Are these countries eating habits similar or different to one another.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
```

Q1: How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
# can use the dim fucntion which outputs 17 rows, and 5 columns
dim(x)
```

```
[1] 17  4
```

```
# preview the first 6 rows
head(x)
```

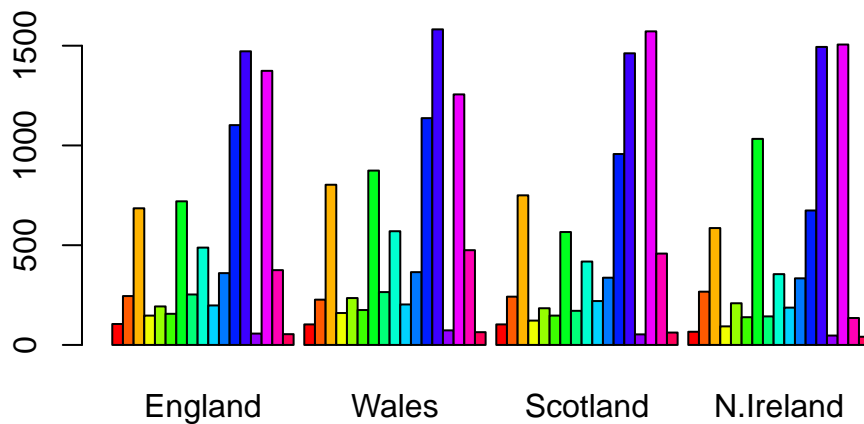
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2: Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

If the method using `x <- x[,-1]` was utilized, and the same cell was run multiple times, it would keep removing rows, which would be problematic, when rendering the document, or when checking your code by rerunning all of your code.

Can visualize food consumption with respect to each country

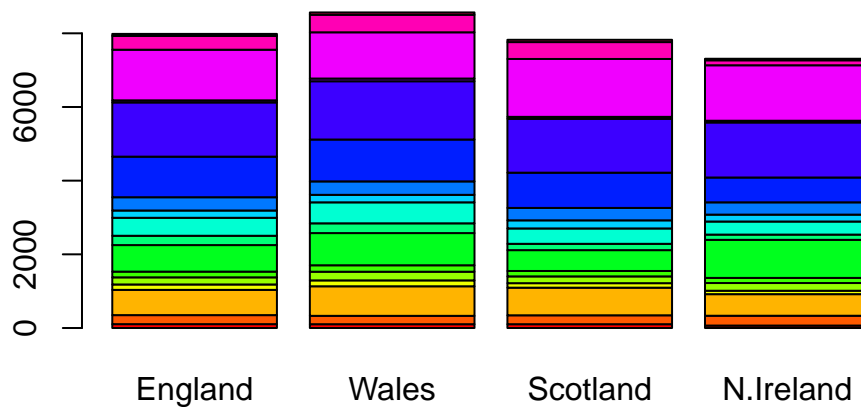
```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

This barplot results in not transposing the data frame, which results in a much worse visualization of the data, since it is much harder to see differences between the variables!

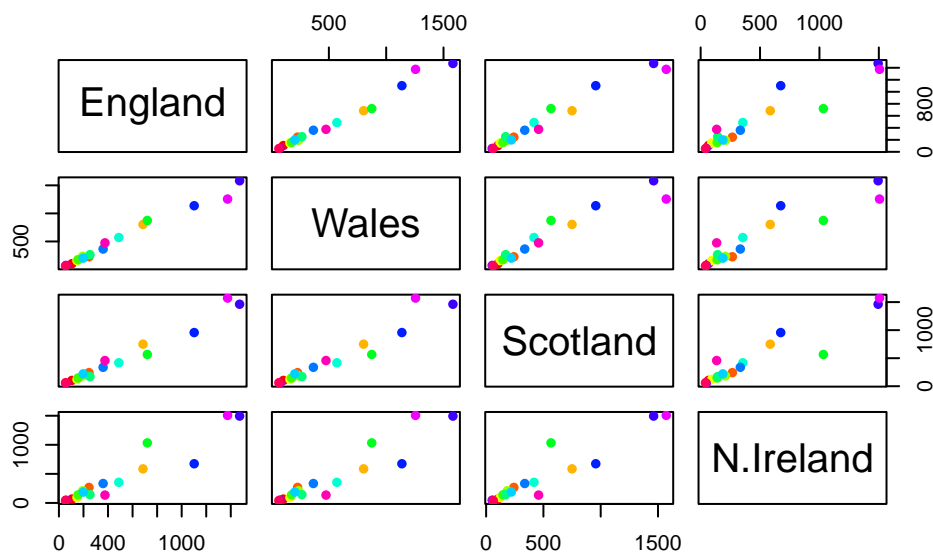
```
# even worse at visualization compared to the other type of barplot
# unable to visualize the difference between variables
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

Pairwise plots compare the different consumption of food type between two different countries, and do it for all of the plots showing the relationship between consumption among each of the countries. For the first row, england is on the y axis for the plots on the first row, the x-axis for each row changes based on country, which allows visualization of consumption between them. If points fall on the same line, it shows similar consumption between the two countries.

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



Q6: What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Looking at the green point among England and Ireland, you can see that there is more consumption in Ireland vs England since it is further right on the X-axis, and there is less consumption in Ireland in comparison to England of the blue point since it is further left on the X-axis

Looking at these types of “pairwise plots” can be helpful but it does not scale well and kind of sucks! There has got to be a better way...

PCA to the rescue!

The main function for PCA in base R is called `prcomp()`. This function will want the food categories to be in the columns, not in the rows, needs us to transpose our input data -i.e. the food categories in columns and the countries as the rows.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00

```
Cumulative Proportion    0.6744    0.9650    1.00000 1.000e+00
```

Let's see what is in our `pca` result object `pca`

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```

```
[1] "prcomp"
```

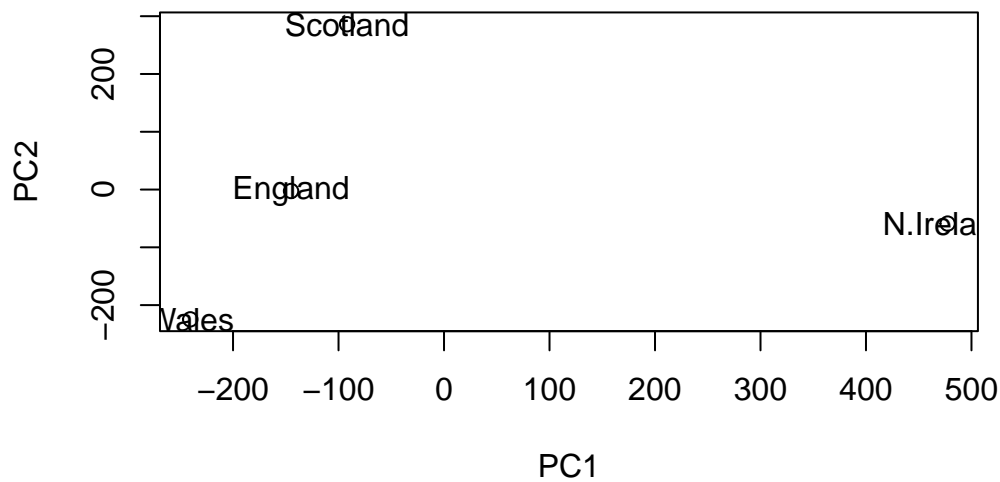
The `pca$x` result object is where we will focus first as this details how the countries are related to each other in terms of our new “axis” (a.k.a. the “PCs”, “eigenvectors”, etc).

```
head(pca$x)
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

Q7: Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2],  
     xlab = 'PC1', ylab = 'PC2')  
  
text(pca$x[,1], pca$x[,2], colnames(x))
```

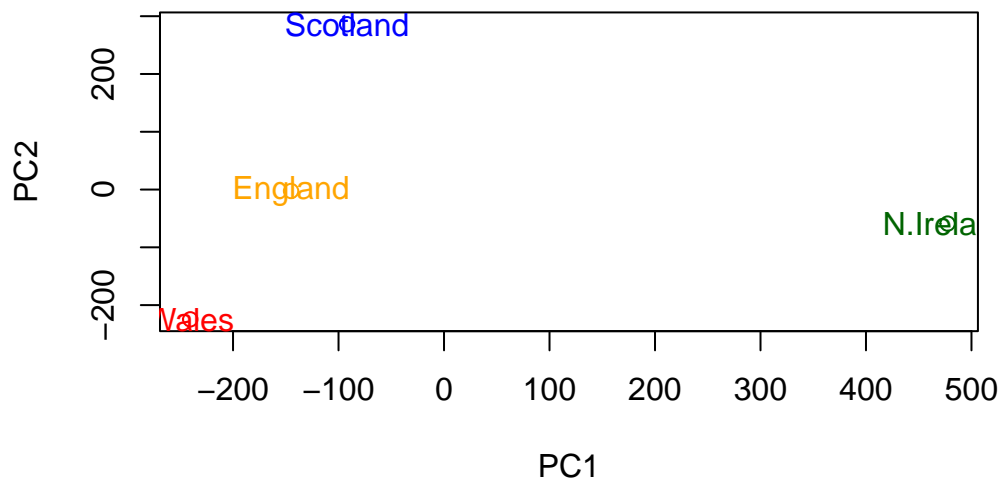


Q8: Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

We will now be plotting PC1 vs PC2

```
plot(pca$x[,1], pca$x[,2], col = c("orange", "red", "blue", "darkgreen"),
     xlab = 'PC1', ylab = 'PC2')

text(pca$x[,1], pca$x[,2], colnames(x), col = c("orange", "red", "blue", "darkgreen"))
```

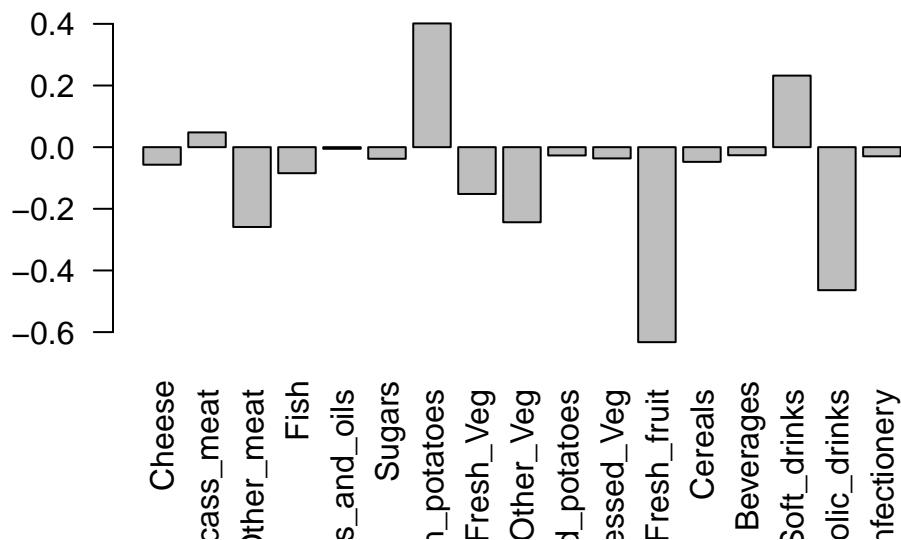


We can look at so called PC loadings result object to see how the original foods contribute to our new PCs (i.e. how the original variables contribute to our new better variables)

```
pca$rotation[,1]
```

Cheese	Carcass_meat	Other_meat	Fish
-0.056955380	0.047927628	-0.258916658	-0.084414983
Fats_and_oils	Sugars	Fresh_potatoes	Fresh_Veg
-0.005193623	-0.037620983	0.401402060	-0.151849942
Other_Veg	Processed_potatoes	Processed_Veg	Fresh_fruit
-0.243593729	-0.026886233	-0.036488269	-0.632640898
Cereals	Beverages	Soft_drinks	Alcoholic_drinks
-0.047702858	-0.026187756	0.232244140	-0.463968168
Confectionery			
-0.029650201			

```
barplot( pca$rotation[,1], las=2 )
```

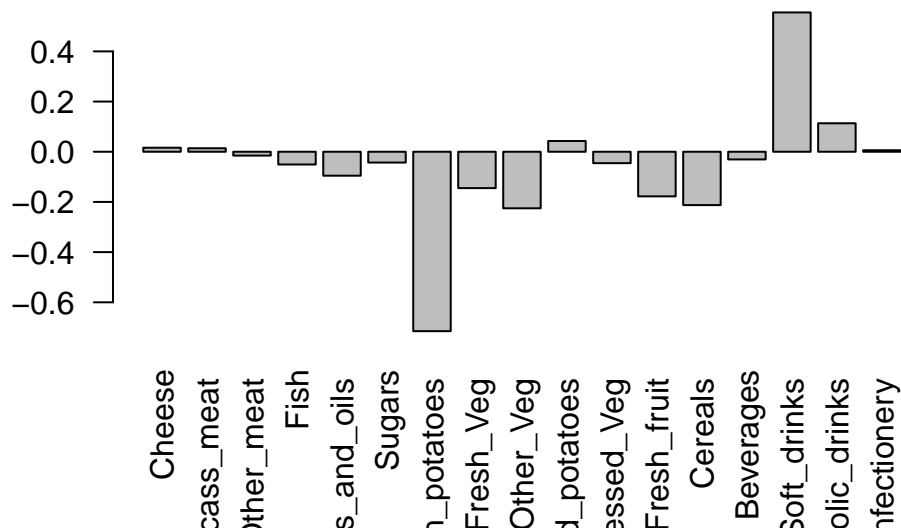


Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
pca$rotation[,2]
```

Cheese	Carcass_meat	Other_meat	Fish
0.016012850	0.013915823	-0.015331138	-0.050754947
Fats_and_oils	Sugars	Fresh_potatoes	Fresh_Veg
-0.095388656	-0.043021699	-0.715017078	-0.144900268
Other_Veg	Processed_potatoes	Processed_Veg	Fresh_fruit
-0.225450923	0.042850761	-0.045451802	-0.177740743
Cereals	Beverages	Soft_drinks	Alcoholic_drinks
-0.212599678	-0.030560542	0.555124311	0.113536523
Confectionery			
0.005949921			

```
barplot( pca$rotation[,2], las=2 )
```

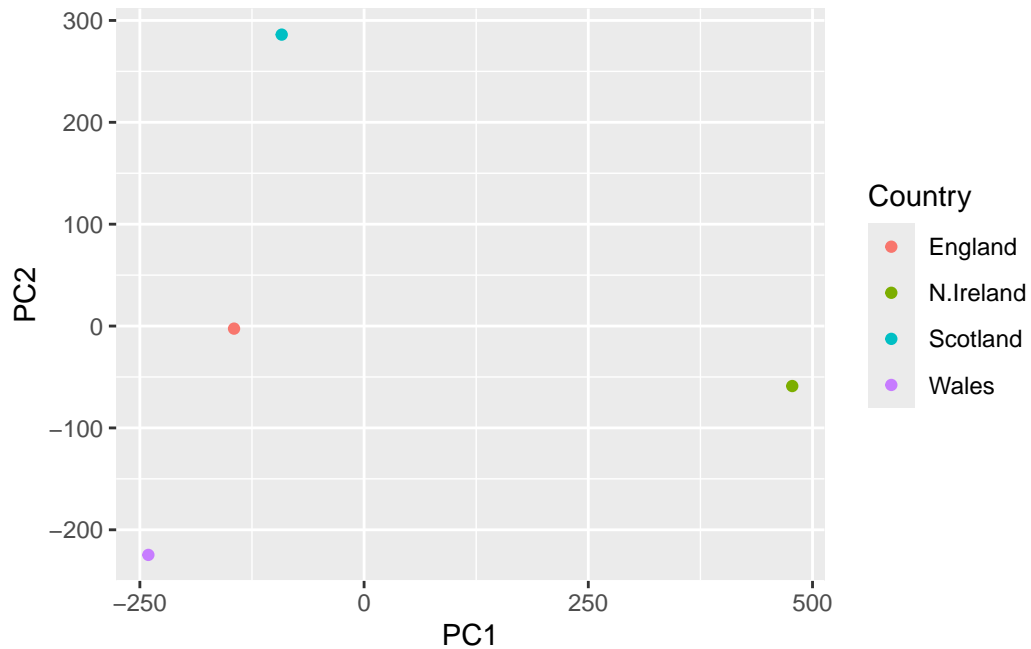
PC2 shows tells us how far up or down the plot the country will be. With a higher positive score, i.e. soft drinks causes Scotland to be further up on the graph. A higher negative score on PC2 shows that the country will be further down on the graph, for example a lower consumption in potatoes causes the country like Whales to be closer to the bottom of the graph.

Using ggplot for the figures:

```
library(ggplot2)

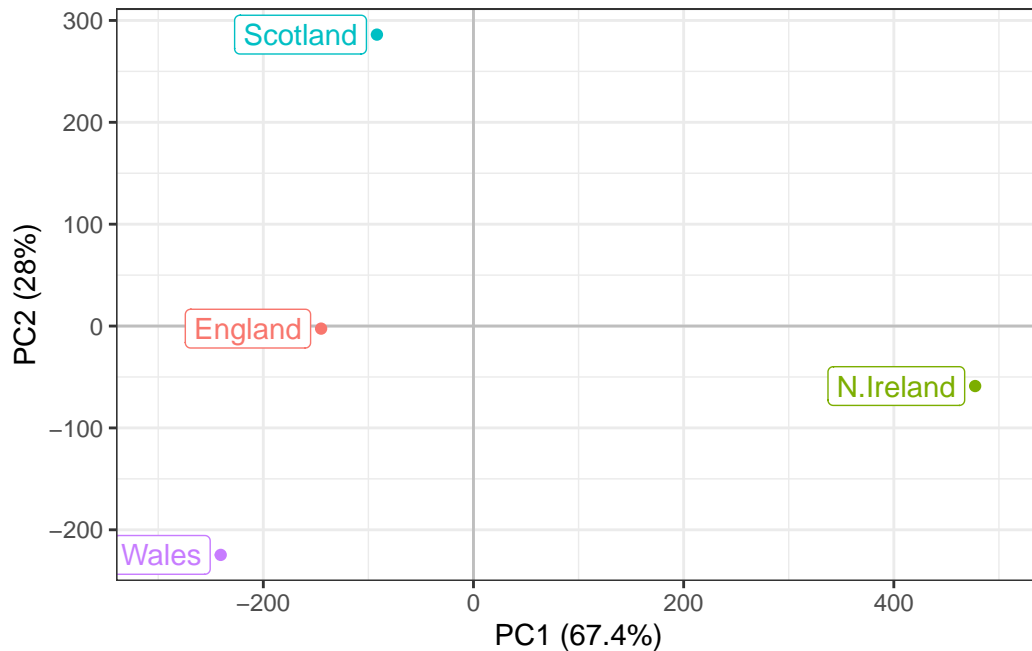
df <- as.data.frame(pca$x)
df_lab <- tibble::rownames_to_column(df, "Country")

# Our first basic plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country) +
  geom_point()
```



making a nicer plot!

```
ggplot(df_lab) +  
  aes(PC1, PC2, col=Country, label=Country) +  
  geom_hline(yintercept = 0, col="gray") +  
  geom_vline(xintercept = 0, col="gray") +  
  geom_point(show.legend = FALSE) +  
  geom_label(hjust=1, nudge_x = -10, show.legend = FALSE) +  
  expand_limits(x = c(-300,500)) +  
  xlab("PC1 (67.4%)") +  
  ylab("PC2 (28%)") +  
  theme_bw()
```



PCA of RNA-seq data:

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10: How many genes and samples are in this data set?

We can determine the number of genes using the `dim()` function which outputs the number of rows and columns, respectively. The number of rows corresponds to the number of genes, which are 100 in this dataset, and the number of columns corresponds to the number of samples which is 10 in this dataset.

```
dim(rna.data)
```

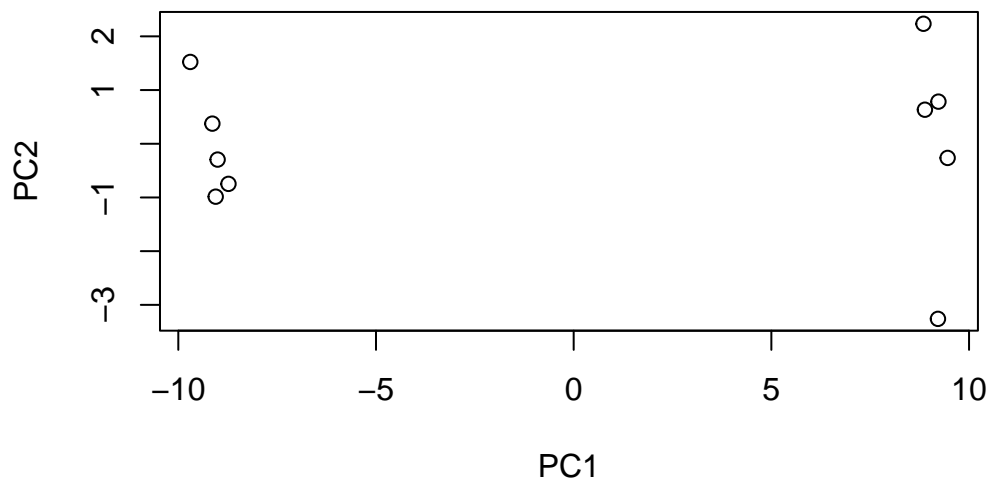
```
[1] 100  10
```

```
## Again we have to take the transpose of our data
```

```
pca1 <- prcomp(t(rna.data), scale=TRUE)
```

```
## Simple un polished plot of pc1 and pc2
```

```
plot(pca1$x[,1], pca1$x[,2], xlab="PC1", ylab="PC2")
```



```
summary(pca1)
```

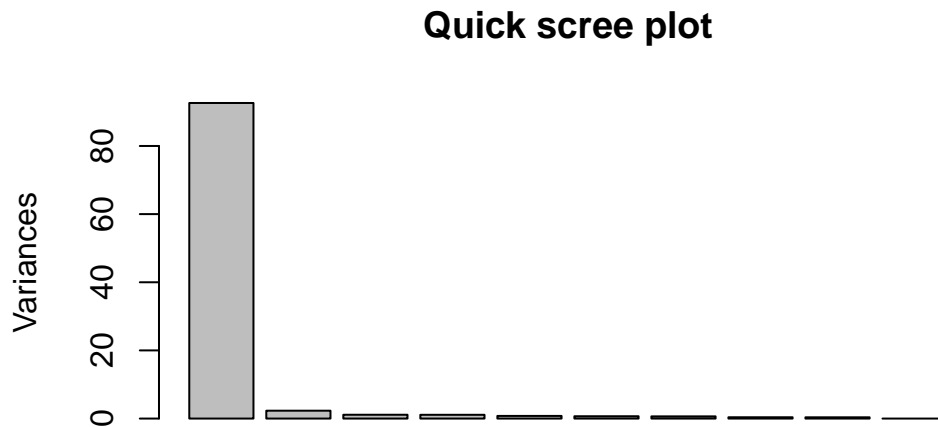
Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251

	PC8	PC9	PC10
Standard deviation	0.62065	0.60342	3.345e-15
Proportion of Variance	0.00385	0.00364	0.000e+00
Cumulative Proportion	0.99636	1.00000	1.000e+00

Make a barplot to quickly visualize this data

```
plot(pca1, main="Quick scree plot")
```



Based on this distribution it looks like there is some skew so we can account for this using `pca1$sdev` which accounts for the standard deviation of the data!

```
## Variance captured per PC
pca1.var <- pca1$sdev^2

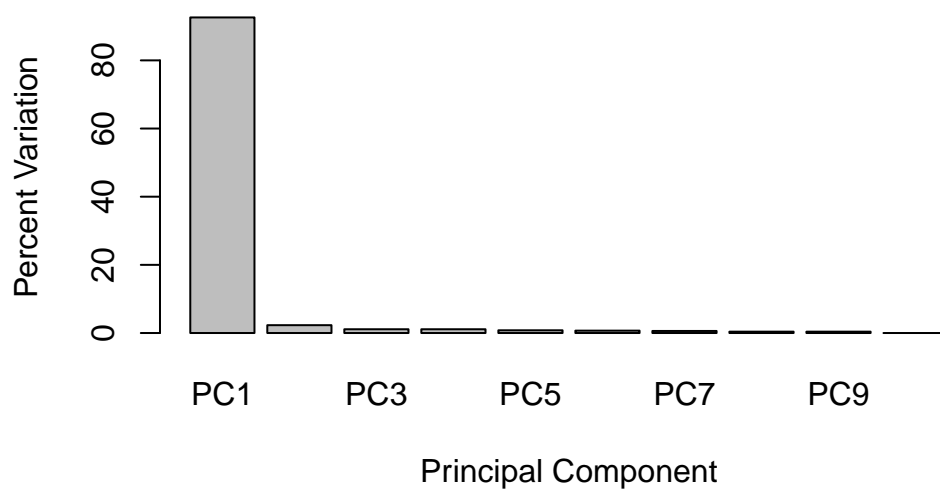
## Percent variance is often more informative to look at
pca1.var.per <- round(pca1.var/sum(pca1.var)*100, 1)
pca1.var.per
```

```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

Using this information we can make another scree plot, to account for the variance!

```
barplot(pca1.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```

Scree Plot

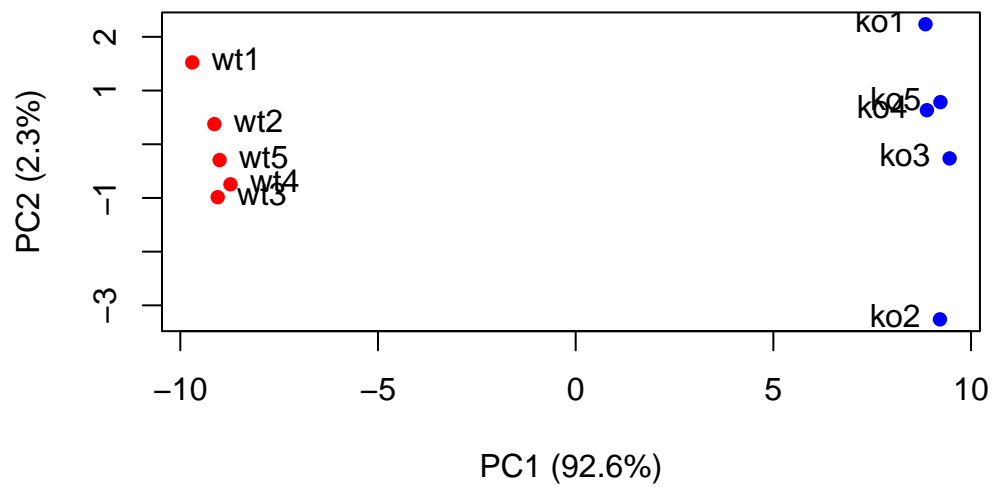


Making a more useful PCA plot:

```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca1$x[,1], pca1$x[,2], col=colvec, pch=16,
      xlab=paste0("PC1 (", pca1.var.per[1], "%)"),
      ylab=paste0("PC2 (", pca1.var.per[2], "%)"))

text(pca1$x[,1], pca1$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```

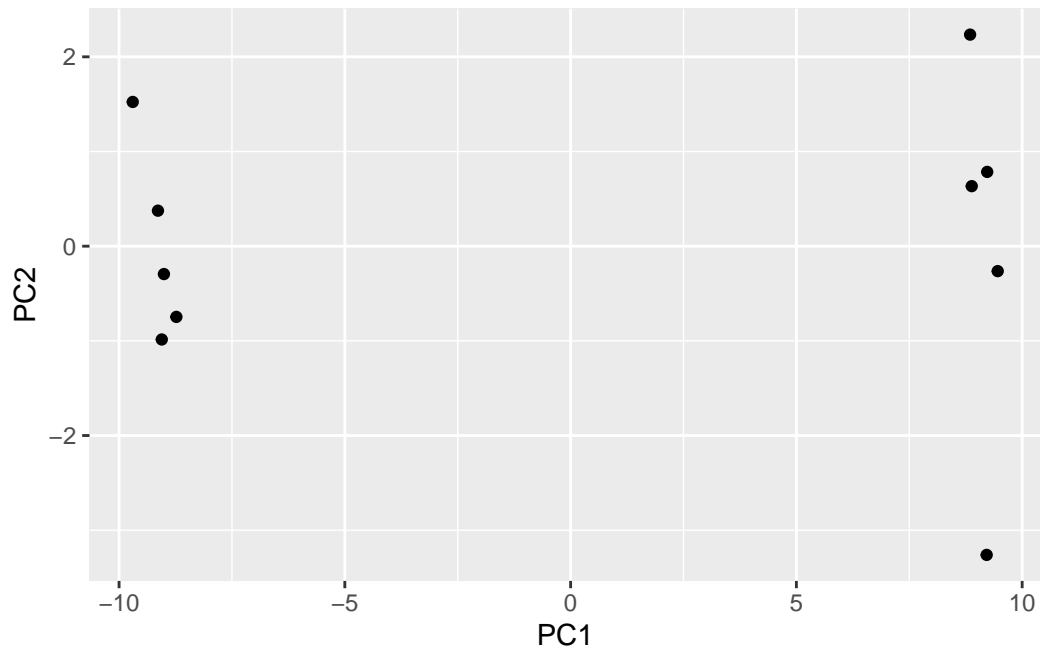


We can also use ggplot to make the figures even more presentable and attractive:

```
library(ggplot2)

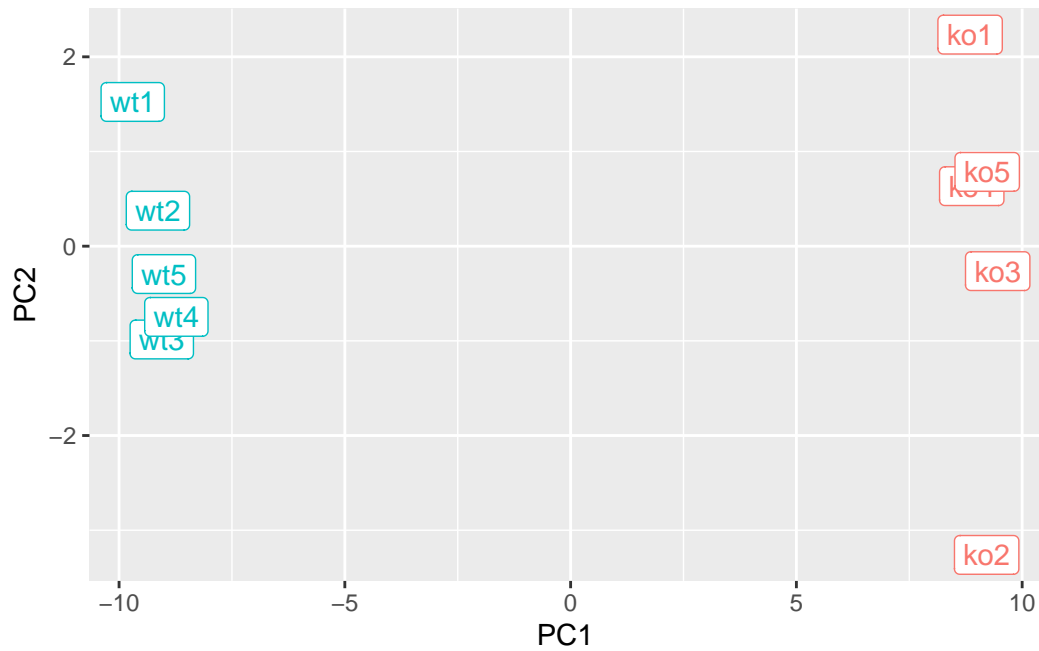
df.rna <- as.data.frame(pca1$x)

# Our first basic plot
ggplot(df.rna) +
  aes(PC1, PC2) +
  geom_point()
```



```
# Add a 'wt' and 'ko' "condition" column
df.rna$samples <- colnames(rna.data)
df.rna$condition <- substr(colnames(rna.data),1,2)

plotrna <- ggplot(df.rna) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
plotrna
```

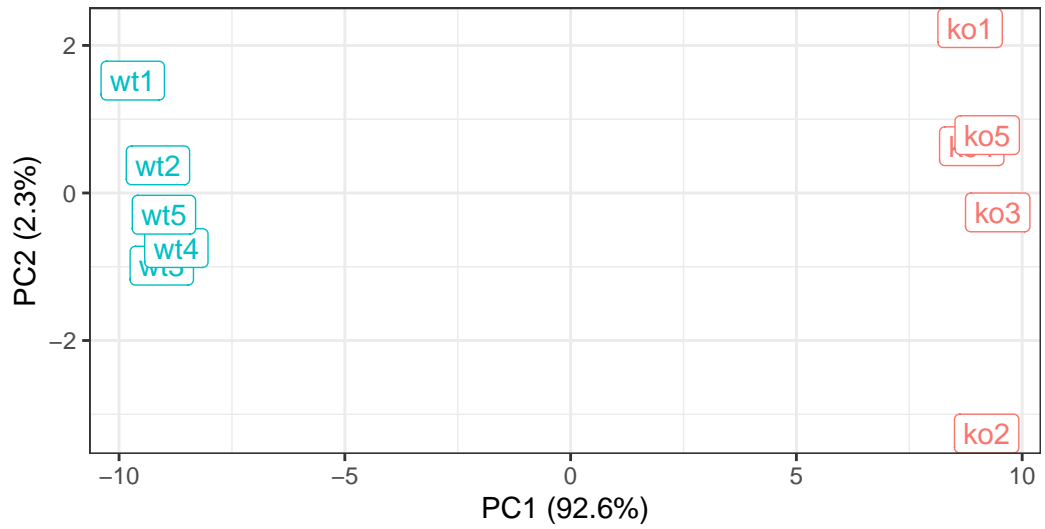



Lastly add final features to polish figure!

```
plotrna + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca1.var.per[1], "%)"),
  y=paste0("PC2 (", pca1.var.per[2], "%)"),
  caption="Class example data") +
theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



Class example data